



Language: EN ▾

How To Set Up Django with Postgres, Nginx, and Gunicorn on Ubuntu 18.04

Posted July 9, 2018 ©190.5k

NGINX

PYTHON

POSTGRESQL

DJANGO

PYTHON FRAMEWORKS

UBUNTU 18.04

DATABASES

By [Justin Ellingwood](#)[Become an author](#)Not using **Ubuntu 18.04**? Choose a different version:

Introduction

Django is a powerful web framework that can help you get your Python application or website off the ground. Django includes a simplified development server for testing your code locally, but for anything even slightly production related, a more secure and powerful web server is required.

In this guide, we will demonstrate how to install and configure some components on Ubuntu 18.04 to support and serve Django applications. We will be setting up a PostgreSQL database instead of using the default SQLite database. We will configure the Gunicorn application server to interface with our applications. We will then set up Nginx to reverse proxy to Gunicorn, giving us access to its security and performance features to serve our apps.

Prerequisites and Goals

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Instance

Sign Up

can learn

how to set this up by running through our [initial server setup guide](#).

We will be installing Django within a virtual environment. Installing Django into an environment specific to your project will allow your projects and their requirements to be handled separately.

Once we have our database and application up and running, we will install and configure the Gunicorn application server. This will serve as an interface to our application, translating client requests from HTTP to Python calls that our application can process. We will then set up Nginx in front of Gunicorn to take advantage of its high performance connection handling mechanisms and its easy-to-implement security features.

Let's get started.

Installing the Packages from the Ubuntu Repositories

To begin the process, we'll download and install all of the items we need from the Ubuntu repositories. We will use the Python package manager `pip` to install additional components a bit later.

We need to update the local `apt` package index and then download and install the packages. The packages we install depend on which version of Python your project will use.

If you are using Django with **Python 3**, type:

```
$ sudo apt update
$ sudo apt install python3-pip python3-dev libpq-dev postgresql postgresql-contrib nginx curl
```

Django 1.11 is the last release of Django that will support Python 2. If you are starting new projects, it is strongly recommended that you choose Python 3. If you still need to use **Python 2**, type:

```
$ sudo apt update
$ sudo apt install python-pip python-dev libpq-dev postgresql postgresql-contrib nginx curl
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

er, the
nginx web

server.

Creating the PostgreSQL Database and User

We're going to jump right in and create a database and database user for our Django application.

By default, Postgres uses an authentication scheme called "peer authentication" for local connections. Basically, this means that if the user's operating system username matches a valid Postgres username, that user can login with no further authentication.

During the Postgres installation, an operating system user named `postgres` was created to correspond to the `postgres` PostgreSQL administrative user. We need to use this user to perform administrative tasks. We can use `sudo` and pass in the username with the `-u` option.

Log into an interactive Postgres session by typing:

```
$ sudo -u postgres psql
```

You will be given a PostgreSQL prompt where we can set up our requirements.

First, create a database for your project:

```
postgres=# CREATE DATABASE myproject;
```

Note: Every Postgres statement must end with a semi-colon, so make sure that your command ends with one if you are experiencing issues.

Next, create a database user for our project. Make sure to select a secure password:

```
postgres=# CREATE USER myprojectuser WITH PASSWORD 'password';
```

Afterwards, we'll modify a few of the connection parameters for the user we just created.

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

Sign Up

We are setting the default encoding to `UTF-8`, which Django expects. We are also setting the default transaction isolation scheme to “read committed”, which blocks reads from uncommitted transactions. Lastly, we are setting the timezone. By default, our Django projects will be set to use `UTC`. These are all recommendations from [the Django project itself](#):

```
postgres=# ALTER ROLE myprojectuser SET client_encoding TO 'utf8';
postgres=# ALTER ROLE myprojectuser SET default_transaction_isolation TO 'read committed';
postgres=# ALTER ROLE myprojectuser SET timezone TO 'UTC';
```

Now, we can give our new user access to administer our new database:

```
postgres=# GRANT ALL PRIVILEGES ON DATABASE myproject TO myprojectuser;
```

When you are finished, exit out of the PostgreSQL prompt by typing:

```
postgres=# \q
```

Postgres is now set up so that Django can connect to and manage its database information.

Creating a Python Virtual Environment for your Project

Now that we have our database, we can begin getting the rest of our project requirements ready. We will be installing our Python requirements within a virtual environment for easier management.

To do this, we first need access to the `virtualenv` command. We can install this with `pip`.

If you are using **Python 3**, upgrade `pip` and install the package by typing:

```
$ sudo -H pip3 install --upgrade pip
$ sudo -H pip3 install virtualenv
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

```
$ sudo -H pip install --upgrade pip
$ sudo -H pip install virtualenv
```

With `virtualenv` installed, we can start forming our project. Create and move into a directory where we can keep our project files:

```
$ mkdir ~/myprojectdir
$ cd ~/myprojectdir
```

Within the project directory, create a Python virtual environment by typing:

```
$ virtualenv myprojectenv
```

This will create a directory called `myprojectenv` within your `myprojectdir` directory. Inside, it will install a local version of Python and a local version of `pip`. We can use this to install and configure an isolated Python environment for our project.

Before we install our project's Python requirements, we need to activate the virtual environment. You can do that by typing:

```
$ source myprojectenv/bin/activate
```

Your prompt should change to indicate that you are now operating within a Python virtual environment. It will look something like this: `(myprojectenv) user@host:~/myprojectdir$`.

With your virtual environment active, install Django, Gunicorn, and the `psycopg2` PostgreSQL adaptor with the local instance of `pip`:

Note: When the virtual environment is activated (when your prompt has `(myprojectenv)` preceding it), use `pip` instead of `pip3`, even if you are using Python 3. The virtual environment's copy of the tool is always named `pip`, regardless of the Python version.

```
(myprojectenv) $ pip install django gunicorn psycopg2-binary
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

Creating and Configuring a New Django Project

With our Python components installed, we can create the actual Django project files.

Creating the Django Project

Since we already have a project directory, we will tell Django to install the files here. It will create a second level directory with the actual code, which is normal, and place a management script in this directory. The key to this is that we are defining the directory explicitly instead of allowing Django to make decisions relative to our current directory:

```
(myprojectenv) $ django-admin.py startproject myproject ~/myprojectdir
```

At this point, your project directory (`~/myprojectdir` in our case) should have the following content:

- `~/myprojectdir/manage.py`: A Django project management script.
- `~/myprojectdir/myproject/`: The Django project package. This should contain the `__init__.py`, `settings.py`, `urls.py`, and `wsgi.py` files.
- `~/myprojectdir/myprojectenv/`: The virtual environment directory we created earlier.

Adjusting the Project Settings

The first thing we should do with our newly created project files is adjust the settings. Open the settings file in your text editor:

```
(myprojectenv) $ nano ~/myprojectdir/myproject/settings.py
```

Start by locating the `ALLOWED_HOSTS` directive. This defines a list of the server's addresses or domain names may be used to connect to the Django instance. Any incoming requests with a **Host** header that is not in this list will raise an exception. Django requires that you set this to prevent a certain class of security vulnerability.

In the square brackets, list the IP addresses or domain names that are associated with your Django server. Each item should be listed in quotations with entries separated by a comma. If you wish requests for an entire domain and any subdomains, prepend a period

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

✕ out

Sign Up

Note: Be sure to include `localhost` as one of the options since we will be proxying connections through a local Nginx instance.

~/myprojectdir/myproject/settings.py

```

. . .
# The simplest case: just add the domain name(s) and IP addresses of your Django server
# ALLOWED_HOSTS = [ 'example.com', '203.0.113.5' ]
# To respond to 'example.com' and any subdomains, start the domain with a dot
# ALLOWED_HOSTS = [ '.example.com', '203.0.113.5' ]
ALLOWED_HOSTS = [ 'your_server_domain_or_IP', 'second_domain_or_IP', . . . , 'localhost' ]

```

Next, find the section that configures database access. It will start with `DATABASES`. The configuration in the file is for a SQLite database. We already created a PostgreSQL database for our project, so we need to adjust the settings.

Change the settings with your PostgreSQL database information. We tell Django to use the `psycopg2` adaptor we installed with `pip`. We need to give the database name, the database username, the database user's password, and then specify that the database is located on the local computer. You can leave the `PORT` setting as an empty string:

~/myprojectdir/myproject/settings.py

```

. . .

DATABASES = {
    'default': {
        'ENGINE': 'django.db.backends.postgresql_psycopg2',
        'NAME': 'myproject',
        'USER': 'myprojectuser',
        'PASSWORD': 'password',
        'HOST': 'localhost',
        'PORT': '',
    }
}

. . .

```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



the static
these

Sign Up

items. The following line tells Django to place them in a directory called `static` in the base project directory:

```
~/myprojectdir/myproject/settings.py
```

```
. . .
```

```
STATIC_URL = '/static/'  
STATIC_ROOT = os.path.join(BASE_DIR, 'static/')
```

Save and close the file when you are finished.

Completing Initial Project Setup

Now, we can migrate the initial database schema to our PostgreSQL database using the management script:

```
(myprojectenv) $ ~/myprojectdir/manage.py makemigrations  
(myprojectenv) $ ~/myprojectdir/manage.py migrate
```

Create an administrative user for the project by typing:

```
(myprojectenv) $ ~/myprojectdir/manage.py createsuperuser
```

You will have to select a username, provide an email address, and choose and confirm a password.

We can collect all of the static content into the directory location we configured by typing:

```
(myprojectenv) $ ~/myprojectdir/manage.py collectstatic
```

You will have to confirm the operation. The static files will then be placed in a directory called `static` within your project directory.

If you followed the initial server setup guide, you should have a UFW firewall protecting

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



o the port

Sign Up

Create an exception for port 8000 by typing:

```
(myprojectenv) $ sudo ufw allow 8000
```

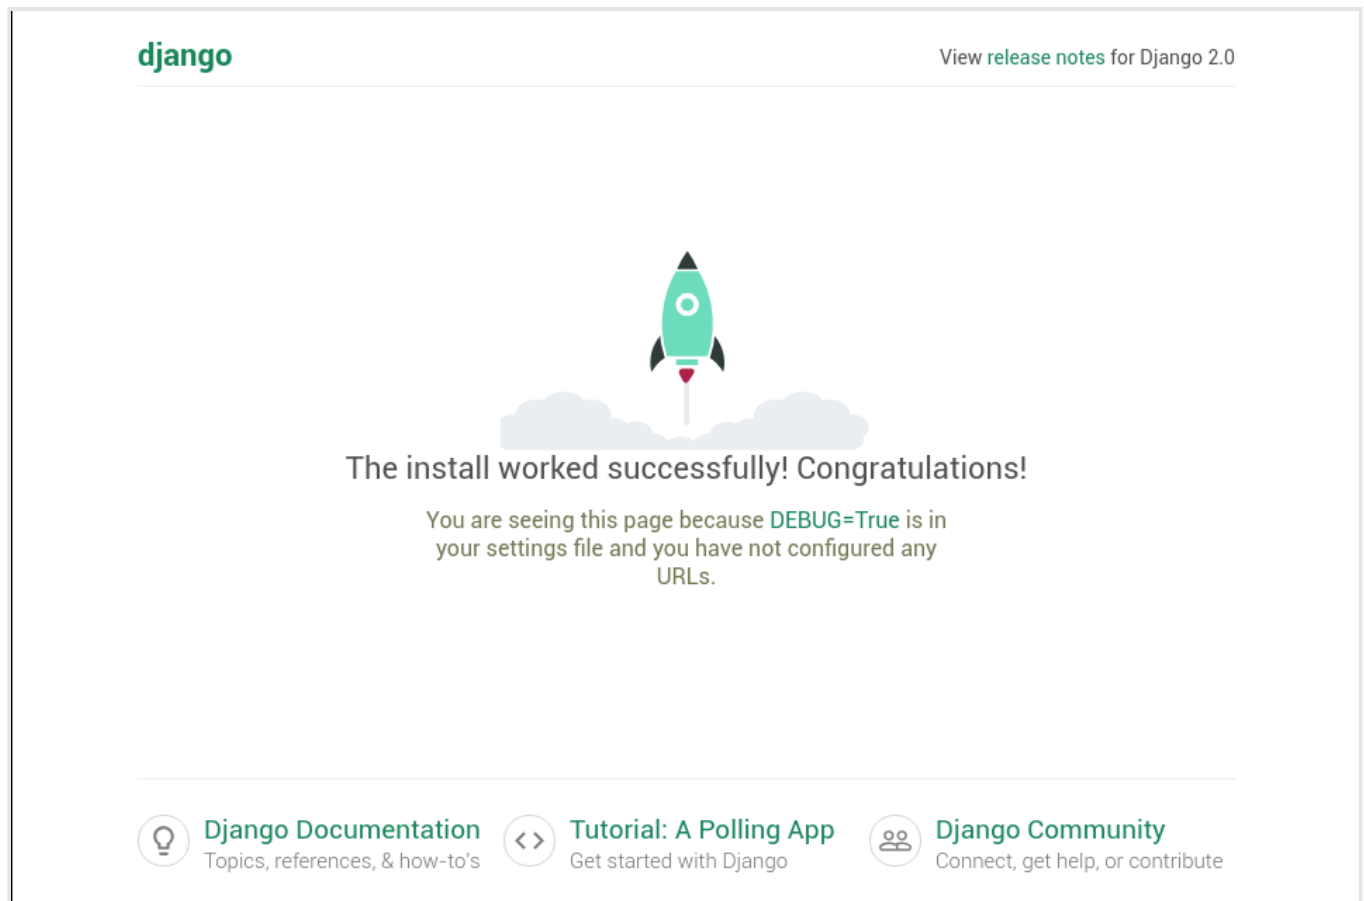
Finally, you can test out your project by starting up the Django development server with this command:

```
(myprojectenv) $ ~/myprojectdir/manage.py runserver 0.0.0.0:8000
```

In your web browser, visit your server's domain name or IP address followed by `:8000`:

```
http://server_domain_or_IP:8000
```

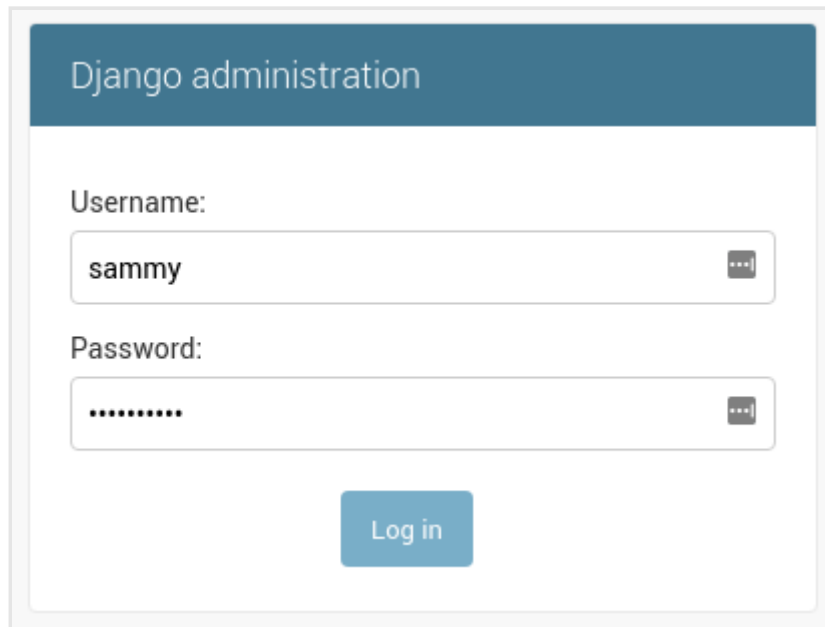
You should see the default Django index page:



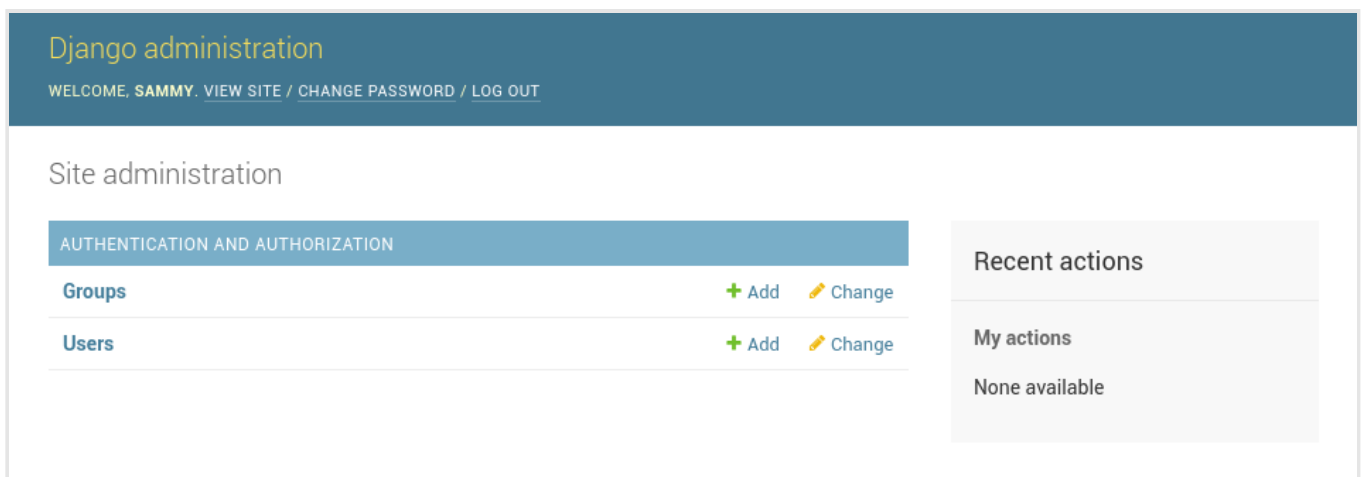
If you append `/admin` to the end of the URL in the address bar, you will be prompted for the administrative username and password you created with the `createsuperuser`

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. ✕

Sign Up

A screenshot of the Django administration login page. It features a dark blue header with the text "Django administration". Below the header, there are two input fields: "Username:" with the value "sammy" and "Password:" with masked characters ".....". To the right of each input field is a small icon for toggling password visibility. At the bottom center is a blue "Log in" button.

After authenticating, you can access the default Django admin interface:

A screenshot of the Django administration dashboard. The header is dark blue with "Django administration" in yellow. Below the header, it says "WELCOME, SAMMY. [VIEW SITE](#) / [CHANGE PASSWORD](#) / [LOG OUT](#)". The main content area is titled "Site administration". On the left, there's a section "AUTHENTICATION AND AUTHORIZATION" with two items: "Groups" and "Users". Each item has a green "+ Add" link and a yellow pencil "Change" link. On the right, there's a "Recent actions" section with a "My actions" subsection that says "None available".

When you are finished exploring, hit **CTRL-C** in the terminal window to shut down the development server.

Testing Gunicorn's Ability to Serve the Project

The last thing we want to do before leaving our virtual environment is test Gunicorn to make sure that it can serve the application. We can do this by entering our project directory and using `gunicorn` to load the project's WSGI module:

```
(myprojectenv) $ cd ~/myprojectdir
(myprojectenv) $ gunicorn --bind 0.0.0.0:8000 myproject.wsgi
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

✕ I've read this

Enter your email address

Sign Up

Note: The admin interface will not have any of the styling applied since Gunicorn does not know how to find the static CSS content responsible for this.

We passed Gunicorn a module by specifying the relative directory path to Django's `wsgi.py` file, which is the entry point to our application, using Python's module syntax. Inside of this file, a function called `application` is defined, which is used to communicate with the application. To learn more about the WSGI specification, click [here](#).

When you are finished testing, hit **CTRL-C** in the terminal window to stop Gunicorn.

We're now finished configuring our Django application. We can back out of our virtual environment by typing:

```
(myprojectenv) $ deactivate
```

The virtual environment indicator in your prompt will be removed.

Creating systemd Socket and Service Files for Gunicorn

We have tested that Gunicorn can interact with our Django application, but we should implement a more robust way of starting and stopping the application server. To accomplish this, we'll make systemd service and socket files.

The Gunicorn socket will be created at boot and will listen for connections. When a connection occurs, systemd will automatically start the Gunicorn process to handle the connection.

Start by creating and opening a systemd socket file for Gunicorn with `sudo` privileges:

```
$ sudo nano /etc/systemd/system/gunicorn.socket
```

Inside, we will create a `[Unit]` section to describe the socket, a `[Socket]` section to define the socket location, and an `[Install]` section to make sure the socket is created at the right time:

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

```
[Unit]
```

```
Description=gunicorn socket
```

```
[Socket]
```

```
ListenStream=/run/gunicorn.sock
```

```
[Install]
```

```
WantedBy=sockets.target
```

Save and close the file when you are finished.

Next, create and open a systemd service file for Gunicorn with `sudo` privileges in your text editor. The service filename should match the socket filename with the exception of the extension:

```
$ sudo nano /etc/systemd/system/gunicorn.service
```

Start with the `[Unit]` section, which is used to specify metadata and dependencies. We'll put a description of our service here and tell the init system to only start this after the networking target has been reached. Because our service relies on the socket from the socket file, we need to include a `Requires` directive to indicate that relationship:

```
/etc/systemd/system/gunicorn.service
```

```
[Unit]
```

```
Description=gunicorn daemon
```

```
Requires=gunicorn.socket
```

```
After=network.target
```

Next, we'll open up the `[Service]` section. We'll specify the user and group that we want to process to run under. We will give our regular user account ownership of the process since it owns all of the relevant files. We'll give group ownership to the `www-data` group so that Nginx can communicate easily with Gunicorn.

We'll then map out the working directory and specify the command to use to start the service. In this case, we'll have to specify the full path to the Gunicorn executable, which is installed within our virtual environment. We will bind the process to the Unix socket we

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

Sign Up

✕ Nginx. We
gunicorn logs.

We can also specify any optional Gunicorn tweaks here. For example, we specified 3 worker processes in this case:

```
/etc/systemd/system/gunicorn.service
```

```
[Unit]
Description=gunicorn daemon
Requires=gunicorn.socket
After=network.target

[Service]
User=sammy
Group=www-data
WorkingDirectory=/home/sammy/myprojectdir
ExecStart=/home/sammy/myprojectdir/myprojectenv/bin/gunicorn \
    --access-logfile - \
    --workers 3 \
    --bind unix:/run/gunicorn.sock \
    myproject.wsgi:application
```

Finally, we'll add an `[Install]` section. This will tell `systemd` what to link this service to if we enable it to start at boot. We want this service to start when the regular multi-user system is up and running:

```
/etc/systemd/system/gunicorn.service
```

```
[Unit]
Description=gunicorn daemon
Requires=gunicorn.socket
After=network.target

[Service]
User=sammy
Group=www-data
WorkingDirectory=/home/sammy/myprojectdir
ExecStart=/home/sammy/myprojectdir/myprojectenv/bin/gunicorn \
    --access-logfile - \
    --workers 3 \
    --bind unix:/run/gunicorn.sock \
    myproject.wsgi:application
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

With that, our `systemd` service file is complete. Save and close it now.

We can now start and enable the Gunicorn socket. This will create the socket file at `/run/gunicorn.sock` now and at boot. When a connection is made to that socket, `systemd` will automatically start the `gunicorn.service` to handle it:

```
$ sudo systemctl start gunicorn.socket
$ sudo systemctl enable gunicorn.socket
```

We can confirm that the operation was successful by checking for the socket file.

Checking for the Gunicorn Socket File

Check the status of the process to find out whether it was able to start:

```
$ sudo systemctl status gunicorn.socket
```

Next, check for the existence of the `gunicorn.sock` file within the `/run` directory:

```
$ file /run/gunicorn.sock
```

Output

```
/run/gunicorn.sock: socket
```

If the `systemctl status` command indicated that an error occurred or if you do not find the `gunicorn.sock` file in the directory, it's an indication that the Gunicorn socket was not able to be created correctly. Check the Gunicorn socket's logs by typing:

```
$ sudo journalctl -u gunicorn.socket
```

Take another look at your `/etc/systemd/system/gunicorn.socket` file to fix any problems before continuing.

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

Currently, if you've only started the `gunicorn.socket` unit, the `gunicorn.service` will not be active yet since the socket has not yet received any connections. You can check this by typing:

```
$ sudo systemctl status gunicorn
```

Output

```
● gunicorn.service - gunicorn daemon
   Loaded: loaded (/etc/systemd/system/gunicorn.service; disabled; vendor preset: enabled)
   Active: inactive (dead)
```

To test the socket activation mechanism, we can send a connection to the socket through `curl` by typing:

```
$ curl --unix-socket /run/gunicorn.sock localhost
```

You should see the HTML output from your application in the terminal. This indicates that Gunicorn was started and was able to serve your Django application. You can verify that the Gunicorn service is running by typing:

```
$ sudo systemctl status gunicorn
```

Output

```
● gunicorn.service - gunicorn daemon
   Loaded: loaded (/etc/systemd/system/gunicorn.service; disabled; vendor preset: enabled)
   Active: active (running) since Mon 2018-07-09 20:00:40 UTC; 4s ago
 Main PID: 1157 (gunicorn)
    Tasks: 4 (limit: 1153)
   CGroup: /system.slice/gunicorn.service
           └─1157 /home/sammy/myprojectdir/myprojectenv/bin/python3 /home/sammy/myprojectdir/myproje
           └─1178 /home/sammy/myprojectdir/myprojectenv/bin/python3 /home/sammy/myprojectdir/myproje
           └─1180 /home/sammy/myprojectdir/myprojectenv/bin/python3 /home/sammy/myprojectdir/myproje
           └─1181 /home/sammy/myprojectdir/myprojectenv/bin/python3 /home/sammy/myprojectdir/myproje
```

```
Jul 09 20:00:40 django1 systemd[1]: Started gunicorn daemon.
```

```
Jul 09 20:00:40 django1 gunicorn[1157]: [2018-07-09 20:00:40 +0000] [1157] [INFO] Starting gunicorn
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

✕ tening at: uni
ng worker: sync
ting worker wit

Sign Up

```
Jul 09 20:00:40 django1 gunicorn[1157]: [2018-07-09 20:00:40 +0000] [1180] [INFO] Booting worker wit
Jul 09 20:00:40 django1 gunicorn[1157]: [2018-07-09 20:00:40 +0000] [1181] [INFO] Booting worker wit
Jul 09 20:00:41 django1 gunicorn[1157]: - - [09/Jul/2018:20:00:41 +0000] "GET / HTTP/1.1" 200 16348
```

If the output from `curl` or the output of `systemctl status` indicates that a problem occurred, check the logs for additional details:

```
$ sudo journalctl -u gunicorn
```

Check your `/etc/systemd/system/gunicorn.service` file for problems. If you make changes to the `/etc/systemd/system/gunicorn.service` file, reload the daemon to reread the service definition and restart the Gunicorn process by typing:

```
$ sudo systemctl daemon-reload
$ sudo systemctl restart gunicorn
```

Make sure you troubleshoot the above issues before continuing.

Configure Nginx to Proxy Pass to Gunicorn

Now that Gunicorn is set up, we need to configure Nginx to pass traffic to the process.

Start by creating and opening a new server block in Nginx's `sites-available` directory:

```
$ sudo nano /etc/nginx/sites-available/myproject
```

Inside, open up a new server block. We will start by specifying that this block should listen on the normal port 80 and that it should respond to our server's domain name or IP address:

```
/etc/nginx/sites-available/myproject
```

```
server {
    listen 80;
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

Next, we will tell Nginx to ignore any problems with finding a favicon. We will also tell it where to find the static assets that we collected in our `~/myprojectdir/static` directory. All of these files have a standard URI prefix of `"/static"`, so we can create a location block to match those requests:

```
/etc/nginx/sites-available/myproject
```

```
server {
    listen 80;
    server_name server_domain_or_IP;

    location = /favicon.ico { access_log off; log_not_found off; }
    location /static/ {
        root /home/sammy/myprojectdir;
    }
}
```

Finally, we'll create a `location / {}` block to match all other requests. Inside of this location, we'll include the standard `proxy_params` file included with the Nginx installation and then we will pass the traffic directly to the Gunicorn socket:

```
/etc/nginx/sites-available/myproject
```

```
server {
    listen 80;
    server_name server_domain_or_IP;

    location = /favicon.ico { access_log off; log_not_found off; }
    location /static/ {
        root /home/sammy/myprojectdir;
    }

    location / {
        include proxy_params;
        proxy_pass http://unix:/run/gunicorn.sock;
    }
}
```

Save and close the file when you are finished. Now, we can enable the file by linking it to the `sites-enabled` directory:

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

Test your Nginx configuration for syntax errors by typing:

```
$ sudo nginx -t
```

If no errors are reported, go ahead and restart Nginx by typing:

```
$ sudo systemctl restart nginx
```

Finally, we need to open up our firewall to normal traffic on port 80. Since we no longer need access to the development server, we can remove the rule to open port 8000 as well:

```
$ sudo ufw delete allow 8000  
$ sudo ufw allow 'Nginx Full'
```

You should now be able to go to your server's domain or IP address to view your application.

Note: After configuring Nginx, the next step should be securing traffic to the server using SSL/TLS. This is important because without it, all information, including passwords are sent over the network in plain text.

If you have a domain name, the easiest way get an SSL certificate to secure your traffic is using Let's Encrypt. Follow [this guide](#) to set up Let's Encrypt with Nginx on Ubuntu 18.04. Follow the procedure using the Nginx server block we created in this guide.

If you do not have a domain name, you can still secure your site for testing and learning with a [self-signed SSL certificate](#). Again, follow the process using the Nginx server block we created in this tutorial.

Troubleshooting Nginx and Gunicorn

If this last step does not show your application, you will need to troubleshoot your

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

Nginx Is Showing the Default Page Instead of the Django Application

If Nginx displays the default page instead of proxying to your application, it usually means that you need to adjust the `server_name` within the `/etc/nginx/sites-available/myproject` file to point to your server's IP address or domain name.

Nginx uses the `server_name` to determine which server block to use to respond to requests. If you are seeing the default Nginx page, it is a sign that Nginx wasn't able to match the request to a sever block explicitly, so it's falling back on the default block defined in `/etc/nginx/sites-available/default`.

The `server_name` in your project's server block must be more specific than the one in the default server block to be selected.

Nginx Is Displaying a 502 Bad Gateway Error Instead of the Django Application

A 502 error indicates that Nginx is unable to successfully proxy the request. A wide range of configuration problems express themselves with a 502 error, so more information is required to troubleshoot properly.

The primary place to look for more information is in Nginx's error logs. Generally, this will tell you what conditions caused problems during the proxying event. Follow the Nginx error logs by typing:

```
$ sudo tail -F /var/log/nginx/error.log
```

Now, make another request in your browser to generate a fresh error (try refreshing the page). You should see a fresh error message written to the log. If you look at the message, it should help you narrow down the problem.

You might see some of the following message:

connect() to unix:/run/gunicorn.sock failed (2: No such file or directory)

This indicates that Nginx was unable to find the `gunicorn.sock` file at the given location. You should compare the `proxy_pass` location defined within

gunicorn.sock file

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

If you cannot find a `gunicorn.sock` file within the `/run` directory, it generally means that the systemd socket file was unable to create it. Go back to the [section on checking for the Gunicorn socket file](#) to step through the troubleshooting steps for Gunicorn.

connect() to unix:/run/gunicorn.sock failed (13: Permission denied)

This indicates that Nginx was unable to connect to the Gunicorn socket because of permissions problems. This can happen when the procedure is followed using the root user instead of a `sudo` user. While systemd is able to create the Gunicorn socket file, Nginx is unable to access it.

This can happen if there are limited permissions at any point between the root directory (`/`) the `gunicorn.sock` file. We can see the permissions and ownership values of the socket file and each of its parent directories by passing the absolute path to our socket file to the `ls -l` command:

```
$ ls -l /run/gunicorn.sock
```

Output

```
f: /run/gunicorn.sock
drwxr-xr-x root root /
drwxr-xr-x root root run
srw-rw-rw- root root gunicorn.sock
```

The output displays the permissions of each of the directory components. By looking at the permissions (first column), owner (second column) and group owner (third column), we can figure out what type of access is allowed to the socket file.

In the above example, the socket file and each of the directories leading up to the socket file have world read and execute permissions (the permissions column for the directories end with `r-x` instead of `---`). The Nginx process should be able to access the socket successfully.

If any of the directories leading up to the socket do not have world read and execute permission, Nginx will not be able to access the socket without allowing world read and execute permissions or making sure group ownership is given to a group that Nginx is a

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

fused”

One message that you may see from Django when attempting to access parts of the application in the web browser is:

```
OperationalError at /admin/login/
could not connect to server: Connection refused
    Is the server running on host "localhost" (127.0.0.1) and accepting
    TCP/IP connections on port 5432?
```

This indicates that Django is unable to connect to the Postgres database. Make sure that the Postgres instance is running by typing:

```
$ sudo systemctl status postgresql
```

If it is not, you can start it and enable it to start automatically at boot (if it is not already configured to do so) by typing:

```
$ sudo systemctl start postgresql
$ sudo systemctl enable postgresql
```

If you are still having issues, make sure the database settings defined in the `~/myprojectdir/myproject/settings.py` file are correct.

Further Troubleshooting

For additional troubleshooting, the logs can help narrow down root causes. Check each of them in turn and look for messages indicating problem areas.

The following logs may be helpful:

- Check the Nginx process logs by typing: `sudo journalctl -u nginx`
- Check the Nginx access logs by typing: `sudo less /var/log/nginx/access.log`
- Check the Nginx error logs by typing: `sudo less /var/log/nginx/error.log`
- Check the Gunicorn application logs by typing: `sudo journalctl -u gunicorn`
- Check the Gunicorn socket logs by typing: `sudo journalctl -u gunicorn.socket`

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

If you update your Django application, you can restart the Gunicorn process to pick up the changes by typing:

```
$ sudo systemctl restart gunicorn
```

If you change Gunicorn socket or service files, reload the daemon and restart the process by typing:

```
$ sudo systemctl daemon-reload  
$ sudo systemctl restart gunicorn.socket gunicorn.service
```

If you change the Nginx server block configuration, test the configuration and then Nginx by typing:

```
$ sudo nginx -t && sudo systemctl restart nginx
```

These commands are helpful for picking up changes as you adjust your configuration.

Conclusion

In this guide, we've set up a Django project in its own virtual environment. We've configured Gunicorn to translate client requests so that Django can handle them. Afterwards, we set up Nginx to act as a reverse proxy to handle client connections and serve the correct project depending on the client request.

Django makes creating projects and applications simple by providing many of the common pieces, allowing you to focus on the unique elements. By leveraging the general tool chain described in this article, you can easily serve the applications you create from a single server.

By [Justin Ellingwood](#)

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

Was this helpful?**Yes****No**[Report an issue](#)

Related

TUTORIAL

How To Use the PDO PHP Extension to Perform MySQL Transactions in PHP on Ubuntu 18.04

A MySQL transaction is a group of logically related SQL commands that are

TUTORIAL

How To Install MariaDB on Ubuntu 18.04

MariaDB is an open-source database management system, commonly used as an alternative for the MySQL portion of the ...

TUTORIAL

How To Create a URL Shortener with Django and GraphQL

GraphQL is an API standard created and open-sourced by Facebook as an alternative

TUTORIAL

How To Create a New Sudo-enabled User on Ubuntu 18.04 [Quickstart]

This guide will show you how to create a new user with `sudo` access on

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.


**Sign Up**

[Ask a question](#)[Search for more help](#)


71 Comments

Leave a comment...

[Sign In to Comment](#)

 [okorieeo](#) July 11, 2018
3 Great tutorial and straight to the point. But how can I host multiple django projects on a single digital ocean droplet and also link them to different domains, for example, if i have two projects e-commerce and booking with domains e-commerce.com and booking.com respectively how can I configure both of them in one droplet? thanks

[Reply](#) [Report](#)

 [jellingwood](#) July 11, 2018
8 [@okorieeo](#) If you want to host multiple projects, you'd repeat a lot of the steps in this guide for your second site. So you'd want to go back and:

1. Create a second database and user for your second site
2. Create a new virtualenv and install Django, gunicorn, and all of the dependencies for your second project in a second location
3. Go through all of the Django steps in the second virtualenv
4. Set up a new Gunicorn systemd service and socket file for the second site

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



[Sign Up](#)

nain and

For the Gunicorn steps, you would probably need to run one instance of gunicorn per site. That means that instead of creating a `gunicorn.service` and `gunicorn.socket` file, you'd want to name them something associated with the site like:

- `e-commerce-gunicorn.service`
- `e-commerce-gunicorn.socket`
- `booking-gunicorn.service`
- `booking-gunicorn.socket`

You'd adjust the socket location referenced in all of the files (the `.service` and `.socket` files as well as the Nginx server block files) to a site-specific location instead of `/run/gunicorn.sock`. So you could use `/run/e-commerce-gunicorn.sock` and `/run/booking-gunicorn.sock` as the socket location for example.

At that point, Nginx should be configured to respond to both of your domains. Depending on which domain is requested, Nginx will select the appropriate server block to use to serve the request (by checking the value of `server_name` in each of the server blocks). Once there, it will forward the dynamic requests to the appropriate Gunicorn instance (based on the socket location specified in the `proxy_pass` directive). That way, once Nginx matches the domain, the requests for the two domains are processed independently.

Hope that helps!

[Reply](#) [Report](#)



[okorieeo](#) July 12, 2018

Thanks! very helpful, but is the port going to be the same for all the projects? if no what are the possible ports to use

[Reply](#) [Report](#)



[jellingwood](#) July 12, 2018

[@okorieeo](#) You shouldn't have to manage ports if you have multiple domain names for your projects. Each domain name can be set up to listen to port 80 (for HTTP traffic) and port 443 (for HTTPS traffic). Nginx will look at the hostname requested by the client to determine which server block to use and the Gunicorn servers don't rely on ports since they use Unix sockets.

[Reply](#) [Report](#)

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

Sign Up

You have `proxypass` as:

`proxypass http://unix:/gunicorn.sock;`

But it should be:

`proxy_pass http://unix:/run/gunicorn.sock;`

Otherwise, great tutorial! I use these write ups all the time so I appreciate you putting the time into making them.

[Reply](#) [Report](#)

 [jellingwood](#) July 11, 2018

1 [@osodeoro](#) Oh, good catch. Thanks for the heads up!

[Reply](#) [Report](#)

 [esteban93](#) July 20, 2018

0 Thanks for the great tutorial. I am able to complete all the steps but when I go check my site it gives me a “502 Bad Gateway” error. When I type:

```
sudo tail -F /var/log/nginx/error.log
```

```
connect() to unix:/home/user/myprojectdirectory/myproject.sock failed
(2: No such file or directory) while connecting to upstream, client: 0.0.0.0,
server 206.189.193.189, request: "Get / HTTP/1.1", upstream:
"http://unix:/home/user/myprojectdirectory/myproject.sock:/", host: "206.189.193.189:80"
```

I already:

- made sure I have the sock file
- double check the `/etc/nginx/sites-available/myproject` : `proxy_pass` was pointing to the right directory and file
- made sure gunicorn runs before nginx

HELP PLEASE!

[Reply](#) [Report](#)

 [AsthanaMe](#) September 3, 2018

0 I am also getting the same error.

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

 ct.

Enter your email address

Sign Up

[Reply](#) [Report](#) **iancarr1412** July 31, 2018

0 Excellent tutorial thanks. I came to it having spent most of a day trying to update the 1-click install to the latest versions of Python and django. This was much more straightforward. Thanks again

[Reply](#) [Report](#) **defcon007** August 8, 2018

1 A note to everyone who is getting failed password authentication error. Postgres automatically changes your entered username to lowercase before saving but when we connect from Django the user is matched case-sensitively.

So, if you created a user “testUser” in psql and enters “testUser” in database connection dictionary of Django you will get a failed password authentication error. Instead, you should write “testuser” and it will work like a charm.

[Reply](#) [Report](#) **tobiasprinz** August 14, 2018

0 Thanks! After reading this the third time and then figuring I had a little typo, I got nearly everything running.

Except for one thing that is absolutely peculiar: I can reach the django admin. I can create new users and new roles with it. I can even list my own models. But once I try to add a new instance of my model, I run into a 500 error.

gunicorn says:

```
gunicorn[2400]: - - [14/Aug/2018:15:41:59 +0000] "GET /admin/mypackage/mymodel/add/ HTTP/1
```



Which I do not understand as getting the base page works

```
gunicorn[2400]: - - [14/Aug/2018:15:33:05 +0000] "GET /admin/mypackage/mymodel/ HTTP/1.0" ;
```



I don't even have a clue how this can be. nginx' hint is even more obscure:

```
[error] 2573#2573: *76 open() "/usr/share/nginx/htmlindex.html" failed (2: No such file or d
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. 

[Sign Up](#)

**tobiasprinz** August 15, 2018

0

Hmm, it was actually a coding error that did not show up on either develop or production, only on this newly created stage system. So that one was specific to my system.

But what might be useful for other people: Even if you see only a plain 500 error page, check whether this is delivered by gunicorn/Django. If it is, you can simply enable debugging in Django and you can stop guessing. Took me a while to realize that this wasn't an nginx error.

[Reply](#) [Report](#)**koryhutchison** August 25, 2018

0

This is a fantastic tutorial. I've already used it twice! Thank you for explaining everything step by step and making it very easy to follow. You're the best!

[Reply](#) [Report](#)**RizkiR** September 9, 2018

0

please answer the question that I made in my personal question session :(

<https://www.digitalocean.com/community/questions/the-browser-shows-this-site-can-t-be-reached>

[Reply](#) [Report](#)**msgturgut** September 14, 2018

0

I followed every step and tried it three times from start but I have the same problem. My project is working fine on my computer but when I deploy it I can only access to admin page most probably because it gets the url from "admin.site.urls" but any other url seems to have a problem. I try to access the other pages but it return "Not Found The requested URL / was not found on this server." . Can anyone help me with this?

[Reply](#) [Report](#)**gajuambi** September 23, 2018

4

Hi, I have been trying from one month, atleast 6 hours a day + 15 hours on weekend. I have deleted and created droplets too on digital ocean to start afresh. nothing seems to work. The farthest i have gotten is to display the nginx page. gunicorn is all well and fine. nginx is all well and fine. no erros from both sides but my site only displays default nginx page. I simply cannot believe that there is no straight forward setup available at an age where we can launch rockets to mars. Not just you but i have followed almost all the tutorials available online in this whole month. There is only on youtube video available on django nginx gunicorn psql and that too is failing. I wish someone just makes a customized image of all

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



or I will pay

[Sign Up](#)

[Reply](#) [Report](#)

^
3 [jamesncl](#) May 5, 2019

I feel your pain. I've also spent literally weeks trying to accomplish something that you would think is really simple: serving my django app on the web. This whole gunicorn and nginx madness is driving me absolutely mental - so many obscure steps and config files to just serve a super-basic django app. Every single tutorial (including this one) I've followed very carefully, step by step, many times, seem to do slightly different things with slightly different commands, and every time I get stuck where something doesn't work, with obscure unhelpful error messages, and it takes days of researching tedious web server admin crap to even guess as to what may possibly be the problem. AAARRGH! #rant

[Reply](#) [Report](#)

^
1 [chefvillas](#) October 28, 2019

The reason nothing ever works is because many of the things that are said in the tutorials are deprecated and don-t work anymore.

For example, upstart command at the nginx server configuration file has been replaced by systemd. You need to create a file /etc/systemd/gunicorn.service, which has a different syntax than the upstart file.

If the tutorial is older than 6 months, forget it, don-t even try it. Most people are following "simple is better than complex" which is 3 years old and it is actually based on another tutorial by a Polish guy from 2013.

So, you will never get to solve it because it is not something you can fix by using your brain. You would have to go through thousands and thousands of pages and configuration files written by thousands of programmers. etc etc

This is why following tutorials is very bad practice. People who have become developers is because they have had a tutor with up to date information and guidance, never by self teaching. Unless you have a good book, with step by step explanations, which is something very difficult to find, because authors don't write to help you but to make a name for themselves.

[Reply](#) [Report](#)

^
0 [chefvillas](#) October 28, 2019

you need to remove the default nginx site

```
sudo rm /etc/nginx/sites-enabled/default
```

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

[Sign Up](#)

^ I followed this tutorial. But I have a little problem here. When I run
 0 `sudo systemctl start gunicorn.socket` it reutrns
 Failed to start gunicorn.socket: Unit gunicorn.socket is not loaded properly: Invalid a

See system logs and '`systemctl status gunicorn.socket`' for details.

And when I run `systemctl status gunicorn.socket` it says `gunicorn.socket - gunicorn daemon`

Loaded: error (Reason: Invalid argument)

Active: inactive (dead)

Here is my gunicorn.socket file

```
[Unit]
Description=gunicorn daemon
Requires=gunicorn.socket
After=network.target

[Service]
User=safaroff
Group=www-data
WorkingDirectory=/home/safaroff/tbc-backend/tbc_bank_proj
ExecStart=/home/safaroff/tbc-backend/tbc_bank_proj/.venv/bin/gunicorn \
    --access-logfile - \
    --workers 3 \
    --bind unix:/run/gunicorn.sock \
    myproject.wsgi:application

[Install]
WantedBy=multi-user.target
```

[Reply](#) [Report](#)

^ [sritanu25](#) November 16, 2018

0 What's the ownership of `/run/gunicorn.sock` ? Is it safaroff ? You might wanna change that to `www-data` and try it out.

[Reply](#) [Report](#)

^ [phillyharper](#) May 10, 2019

I'm getting the same error - I also don't have a file in `/run` called `gunicorn.sock`. First

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. ✕

Enter your email address

Sign Up

**bengaltiger** December 14, 2019

0

Hello Juniorbee, I see this in your file “myproject.wsgi” but myproject should be your project name. Also you referred this pasted content is gunicorn.socket when this pasted content is suppose to be gunicorn.service, isn't it?

[Reply](#) [Report](#)**mario535974** October 1, 2018

1

Great Tutorial, I'm in the middle of it. Just one detail, when creating the virtual env in python3 we must use the command: `virtualenv -p python3 envname`

[Reply](#) [Report](#)**bipoza** October 12, 2018

0

Thank you! You have helped me with my problem. I'm using it with Django 2

[Reply](#) [Report](#)**DripDropDev** October 3, 2018

0

I have followed the tutorial and everything is working. How would I add version control to this?

p.s. Thank you for this tutorial!

[Reply](#) [Report](#)**john971** October 5, 2018

0

Is there any “real world” example of how to do something similar like this with docker? I've succeeded on this one but I'd prefer to dockerize my Django application.

[Reply](#) [Report](#)**bradenholt** October 10, 2018

0

If you're getting a Bad Gateway, as I was, here's a couple things to check

1) Make sure to add localhost to your `allowed_hosts`

2) If you think fixed some things and it's still not working, try resetting your cache

[Reply](#) [Report](#)**bukaloenis** October 19, 2018

1

Great tutorial however, I followed everything step by step and I cannot get past the Gunicorn section. I keep getting, no module named '<my project name without brackets>.wsgi'.

[Reply](#) [Report](#)

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.

**Sign Up**

^ [potterdavidm](#) November 3, 2018

1

I had this error.

Had to add one more directory to the WorkingDirectory in .service file

WorkingDirectory = /home/sammy/myprojectdirectory/mydjango project

[Reply](#) [Report](#)

^ [jamesncl](#) May 8, 2019

0

Yep, there's a mistake in the tutorial. Seriously, i wish these tutorials were actually checked before they are posted!

[Reply](#) [Report](#)

^ [ehmatthes](#) October 20, 2018

0 Thanks for such a clear tutorial, it's part of what keeps me on DO. I ran into an issue at this step:

```
curl --unix-socket /run/gunicorn.sock localhost
```

I have a few Django settings stored in environment variables, in a file called **.env**. The environment variables were not getting loaded, so I saw an `ImproperlyConfigured` error in the logs. The solution was straightforward. I modified the `gunicorn.service` file to include an `EnvironmentFile` setting:

```
[Service]
User=sammy
Group=www-data
WorkingDirectory=/path/to/project
EnvironmentFile=/path/to/.env
ExecStart=...
```

When I looked up how to load environment variables from systemd, I saw references back to this tutorial. Other readers are running into this issue, so I thought I'd include a note about it here.

[Reply](#) [Report](#)

^ [levicruz](#) October 29, 2018

0

hello, great tutorial!!

I follow all steps,

but. now i receive this error:

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up


```
Request Method: GET
Request URL:      http://68.183.27.123/
Django Version: 2.1.2
Exception Type: TemplateDoesNotExist
Exception Value:
home/home.html
Exception Location: /home/levi/sis_ecovia_dir/ecoviaenv/lib/python3.6/site-packages/django/
Python Executable: /home/levi/sis_ecovia_dir/ecoviaenv/bin/python3
Python Version: 3.6.7
Python Path:
['/home/levi/sis_ecovia_dir',
 '/home/levi/sis_ecovia_dir/ecoviaenv/bin',
 '/home/levi/sis_ecovia_dir/ecoviaenv/lib/python3.6.zip',
 '/home/levi/sis_ecovia_dir/ecoviaenv/lib/python3.6',
 '/home/levi/sis_ecovia_dir/ecoviaenv/lib/python3.6/lib-dynload',
 '/usr/lib/python3.6',
 '/home/levi/sis_ecovia_dir/ecoviaenv/lib/python3.6/site-packages']
Server time:      Mon, 29 Oct 2018 14:09:05 +0000
```

but in localhost it's worker;

[Reply](#) [Report](#)

Load More Comments



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.



Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics. ✕

Sign Up

BECOME A CONTRIBUTOR

You get paid; we donate to tech nonprofits.

**GET OUR BIWEEKLY NEWSLETTER**

Sign up for Infrastructure as a Newsletter.

**COVID-19 SUPPORT PROGRAM**

Working on something related to COVID-19? DigitalOcean would like to help.

[Featured on Community](#) [Kubernetes Course](#) [Learn Python 3](#) [Machine Learning in Python](#)
[Getting started with Go](#) [Intro to Kubernetes](#)

[DigitalOcean Products](#) [Droplets](#) [Managed Databases](#) [Managed Kubernetes](#) [Spaces Object Storage](#)
[Marketplace](#)

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

DigitalOcean makes it simple to launch in the cloud and scale up as you grow – whether you're running one virtual machine or ten thousand.

Learn More



© 2020 DigitalOcean, LLC. All rights reserved.

Company

[About](#)
[Leadership](#)
[Blog](#)
[Careers](#)
[Partners](#)
[Referral Program](#)
[Press](#)
[Legal & Security](#)

Products

[Products Overview](#)
[Pricing](#)
[Droplets](#)
[Kubernetes](#)
[Managed Databases](#)
[Spaces](#)
[Marketplace](#)
[Load Balancers](#)
[Block Storage](#)
[Tools & Integrations](#)

[API](#)
[Documentation](#)
[Release Notes](#)

Community

[Tutorials](#)
[Q&A](#)

Contact

[Get Support](#)
[Trouble Signing In?](#)

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Sign Up

Meetups

Write for DOnations

Droplets for Demos

Hatch Startup Program

Shop Swag

Research Program

Open Source

Code of Conduct

Sign up for our newsletter. Get the latest tutorials on SysAdmin and open source topics.



Enter your email address

Sign Up